AI Agent Code Coverage: Ablation Study

Prompt, Knowledge, Model, and Framework Effects on AI-Generated Tests

Mark Pollack

March 2026 — DevNexus

1 Setup

Task: Given a Spring Boot project with production code and *no tests*, generate a JUnit test suite targeting ≥80% line coverage.

Dataset: 5 Spring Getting Started guides (trivial, 2–3 classes) + Spring Pet Clinic (complex, 15+ classes, 4 layers).

Notation: `Agent:Model` — e.g., `CC:Sonnet` = Claude Code CLI running Claude Sonnet 4.6; `Loopy:Haiku` = Loopy framework (Spring AI agent) running Claude Haiku 4.5.

Variants tested — two axes of variation:

Agent:Model	Prompt	Knowledge	What it tests
CC:Sonnet (naive)	v0-naive	none	Baseline — minimal instructions
CC:Sonnet (hardened)	v1-hardened	none	Does prompt quality matter?
CC:Sonnet (+3 KB)	v2-with-kb	3 files	Does curated knowledge help?
CC:Sonnet (+full KB)	v2-with-kb	full KB	Does more knowledge help more?
CC:Sonnet (two-phase)	v3-explore+act	full KB	Does structured execution help?
CC:Sonnet (+SAE)	v1-hardened+SAE	none	Does pre-computed analysis help Sonnet?
CC:Haiku	v1-hardened	none	Can a cheap model match Sonnet?
CC:Haiku (+SAE)	v1-hardened+SAE	none	Does pre-computed analysis help Haiku?
Loopy:Haiku	v1-hardened	none	Does the framework matter?
Loopy:Sonnet	v1-hardened	none	Does Sonnet fix Loopy's quality gap?
Loopy:Qwen3	v1-hardened	none	Can a free local model do the job?

Scoring: 4-tier automated judge cascade — build success (T0), coverage preservation (T1), coverage improvement + golden test similarity (T2), LLM practice adherence across 6 criteria (T3).

Table 1: Variant comparison (N=1 per cell, sorted by variant progression)

Agent:Model	N	Pass%	Cov%	T3	Eff	Total	Per Item
CC:Sonnet (naive)	5	100%	95.6	0.62	0.88	$4.57	$0.91
CC:Sonnet (hardened)	5	100%	85.7	0.80	0.94	$4.17	$0.83
CC:Sonnet (+3 KB)	5	100%	95.6	0.70	0.84	$4.98	$1.00
CC:Sonnet (+full KB)	5	100%	84.9	0.75	0.85	$5.13	$1.03
CC:Sonnet (two-phase)	6	100%	94.3	0.72	0.65	$16.48	$2.75
CC:Sonnet (forge)	1	100%	95.3	0.95	0.17	$8.55	$8.55
CC:Sonnet (+SAE)	1	100%	71.4	0.93	0.86	$0.92	$0.92
CC:Haiku	6	17%	78.5	0.44	0.77	$2.88	$0.48
CC:Haiku (+SAE)	1	0%	71.4	0.35	0.41	$0.48	$0.48
Loopy:Sonnet	1	100%	94.6	0.73	0.86	$2.72	$2.72
Loopy:Haiku	6	33%	71.4	0.49	0.86	$15.30	$2.55
Loopy:Qwen3	1	0%	0	—	0.73	$5.01	$5.01

Table 2: Cost and token breakdown per variant

Agent:Model	Total	Per Item	In Tok	Out Tok	Think Tok	Duration
CC:Sonnet (naive)	$4.57	$0.91	109	55,430	10,160	1513s
CC:Sonnet (hardened)	$4.17	$0.83	137	49,343	7,791	1481s
CC:Sonnet (+3 KB)	$4.98	$1.00	137	59,472	11,328	1576s
CC:Sonnet (+full KB)	$5.13	$1.03	1,119	56,017	9,924	1671s
CC:Sonnet (two-phase)	$16.48	$2.75	6,953	158,614	29,197	3046s
CC:Sonnet (forge)	$8.55	$8.55	96	60,528	6,854	1435s
CC:Sonnet (+SAE)	$0.92	$0.92	33	8,747	1,462	391s
CC:Haiku	$2.88	$0.48	1,905	152,835	22,509	2271s
CC:Haiku (+SAE)	$0.48	$0.48	317	25,680	3,676	370s
Loopy:Sonnet	$2.72	$2.72	854,353	10,237	0	314s
Loopy:Haiku	$15.30	$2.55	18,336,594	158,691	0	2470s
Loopy:Qwen3	$5.01	$5.01	1,547,477	24,486	0	1421s

2 Results at a Glance

3 Key Findings

3.1 Finding 1: Prompt Engineering Is the Biggest Lever

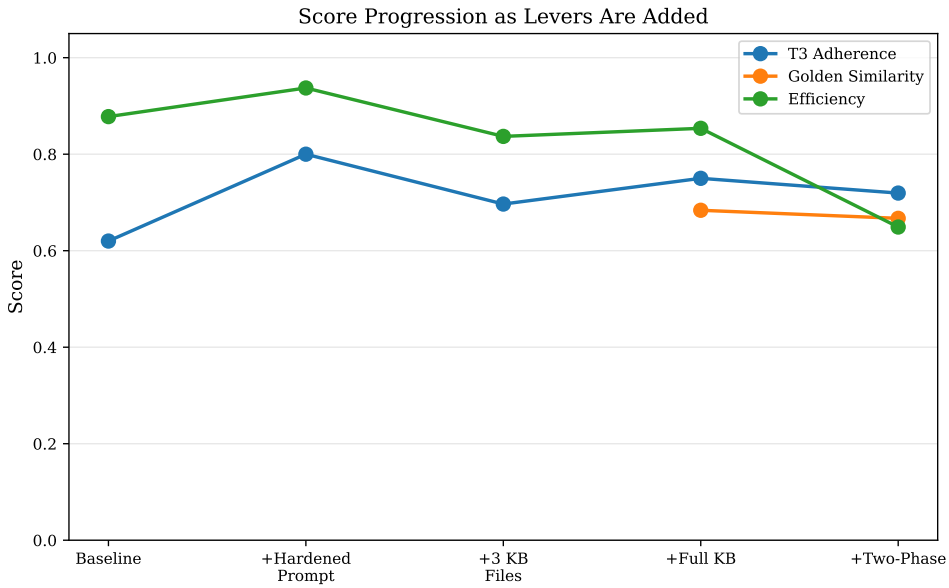

Figure 1: Score progression as levers are added. The jump from naive to hardened prompt is the dominant improvement.

CC:Sonnet (hardened) delivers the best quality: T3=0.81, 100% pass rate, $1.39/item. Knowledge injection and two-phase execution add cost without measurable benefit on simple targets.

3.2 Finding 2: Model Is the Quality Gate, Framework Is the Cost Gate

Quality depends on model capability. Cost depends on framework engineering. These are **independent axes**.

Agent:Model	Framework	T3	Cost	Input Tokens	Passed
CC:Haiku	Claude Code	0.40	$0.45	300	No
Loopy:Haiku	Loopy (no compaction)	0.49	$2.55	18,336,594	No
CC:Sonnet	Claude Code	0.70	$0.92	—	Yes
Loopy:Sonnet	Loopy (with compaction)	0.73	$2.72	854,353	Yes

Table 3: gs-accessing-data-jpa: same item, same prompt across four agent/model combinations.

Read the rows: Haiku fails regardless of framework (T3 0.40–0.49, both fail jury). Sonnet passes regardless of framework (T3 0.70–0.73, both pass jury). The model determines *whether it works*. The framework determines *what it costs*.

Read the columns: Loopy:Haiku without compaction: 18M input tokens. Loopy:Sonnet with compaction (threshold 0.3): 854K input tokens — a **21**× reduction. Context compaction closes most of the framework gap, but Loopy:Sonnet still costs **3**× more than CC:Sonnet at API rates.

First Loopy:Sonnet run hit the $5 cost cap at turn 54 with compaction threshold 0.5. Lowering to 0.3 (compact at 60K tokens) let it complete at $2.72. Compaction tuning matters —

too late and you burn budget on stale context.

3.3 Finding 3: There Is a Model Floor

Both CC:Haiku (T3=0.44) and Loopy:Haiku (T3=0.49) achieve decent **coverage** (71–94%) but fail the quality jury. Common failures: wrong test annotations (`@SpringBootTest` everywhere instead of `@WebMvcTest`), wrong injection patterns, no edge case tests. Below the model floor, no amount of framework engineering, knowledge injection, or prompt tuning helps. Haiku and Qwen3 are below that floor for Spring testing.

3.4 Finding 4: Knowledge Amplifies Capability

We built a lightweight **Stable Analysis Envelope (SAE)** — a pre-computed project report generated after compilation that gives the agent: dependency inventory, component classification (controllers, services, repos, entities), and **copy-paste import blocks** with full package names for each test slice.

We iterated the SAE through two versions, then tested each against two models:

SAE v1 (short annotations)
Component roles and dependency list only. Tells the agent *what exists* but not *how to import it*. Result: the agent still launched `jar tf` commands to discover package names.

SAE v2 (full import blocks)
Adds ready-to-paste import statements for every test slice (`@WebMvcTest`, `@DataJpaTest`, etc.). Result: zero jar exploration — the agent writes tests immediately.

Model	SAE	T3	Cov%	Cost	jar-inspect	Passed
Haiku	none	0.43	71.4	$0.44	0	No
Haiku	v1	0.45	71.4	$0.39	8	No
Haiku	v2	0.35	71.4	$0.48	0	No
Sonnet	none	0.81	85.7	$0.83	—	Yes
Sonnet	v2	0.93	71.4	$0.92	0	Yes

Table 4: SAE ablation on gs-rest-service (single item, N=1). Rows are grouped by model. Haiku with any SAE version still fails the quality jury. Sonnet with SAE v2: T3=0.93, the highest single-item score in the experiment.

Three things stand out:

1. **SAE v2 eliminates exploratory waste.** Zero `jar tf` calls — the agent writes tests immediately instead of probing JARs.
2. **SAE amplifies capable models.** Sonnet + SAE: T3 jumps from 0.81 to 0.93. Same model, same prompt, just better context.
3. **SAE cannot lift weak models.** Haiku + SAE v2: T3 *drops* to 0.35. The information is all there; Haiku can't use it correctly.

Implication: pre-computed analysis is both an efficiency lever (fewer wasted tool calls) and a quality lever — but only when paired with a model above the capability floor. Knowledge amplifies capability; it cannot create it.

3.5 Finding 5: Local Models Are Not Ready

Loopy:Qwen3 (Qwen3-Coder 30B via LM Studio) scored T3=0.30 with 0% pass rate. Even with 300 turns and 45-minute timeout, it couldn't write a single working test for the simplest project. It spent turns reading files and writing todo lists instead of writing tests.

3.6 Finding 6: Partial Knowledge Costs More Than Full or None

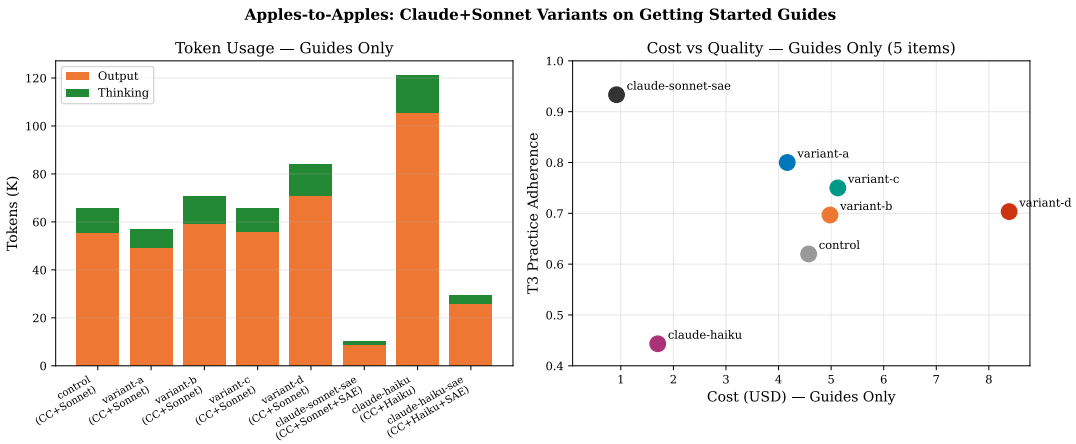

Figure 2: Guides-only comparison (apples-to-apples, 5 items). CC:Sonnet (+3 KB) uses more tokens than (+full KB) — partial knowledge triggers exploratory `jar tf` inspection of Spring Boot JARs to fill gaps.

3.7 Finding 7: Two-Phase Costs 2× for No Gain

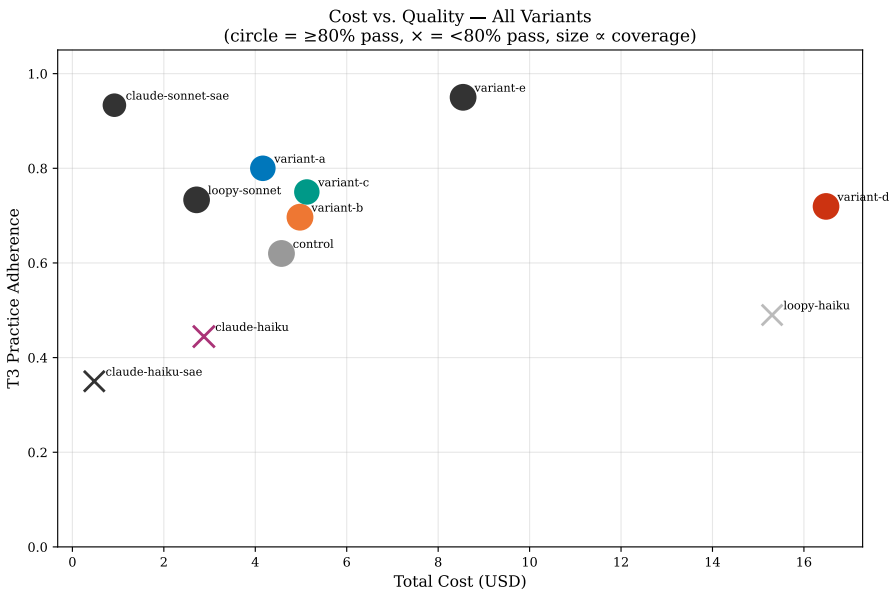

Figure 3: Cost vs. quality across all variants. CC:Sonnet (+SAE) is the new efficient frontier (T3=0.93, \$0.92). ○ = ≥80% pass, × = <80% pass, size ∝ coverage.

4 Agent Behavior

4.1 Token Consumption

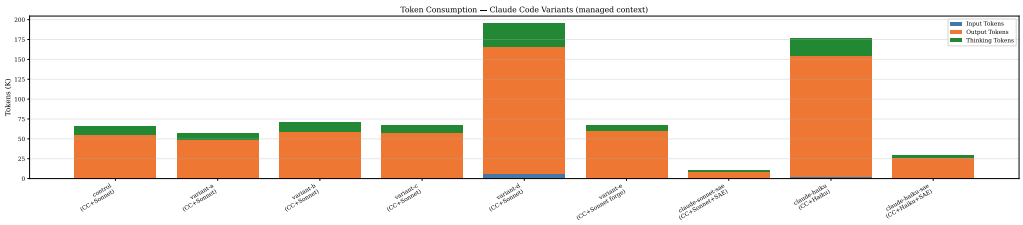

Figure 4: Token consumption — Claude Code variants. Context is managed by the CLI (compaction, summarization). Two-phase variant has the highest thinking tokens.

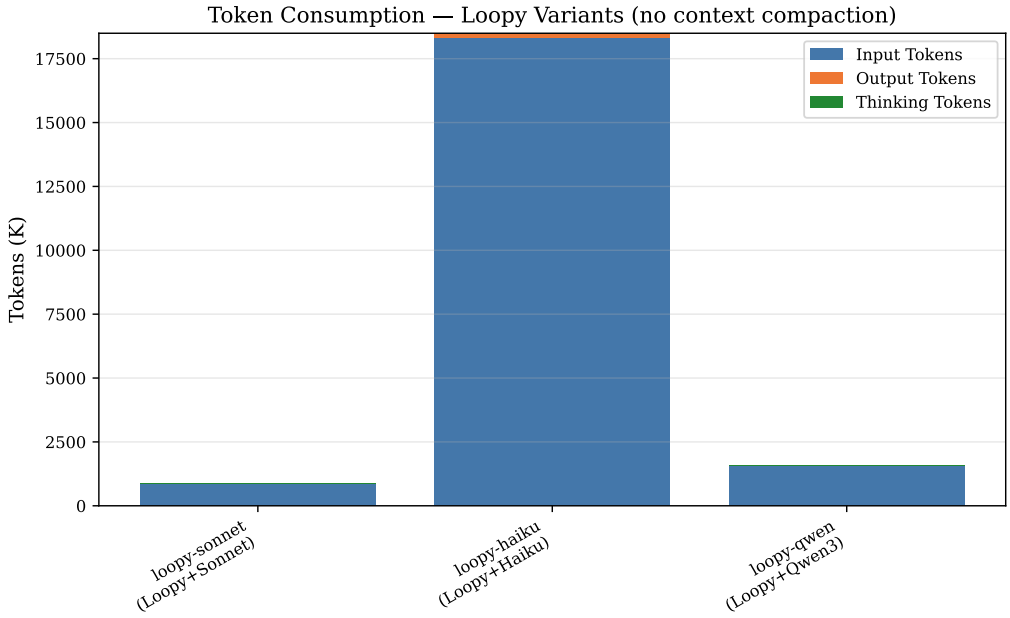

Figure 5: Token consumption — Loopy variants (note different Y scale). Without context compaction, input tokens grow monotonically — every prior tool result is re-sent on each turn. This is the dominant cost driver.

4.2 Tool Usage Patterns

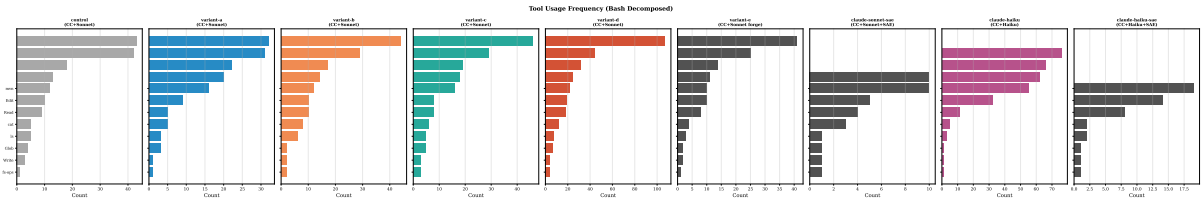

Figure 6: Tool call frequency per variant (bash decomposed to actual utilities). `jar-inspect` = `jar tf` commands probing Spring Boot JARs in `~/.m2` to discover test annotations. Knowledge-injected variants need fewer jar inspections.

4.3 Execution Timelines

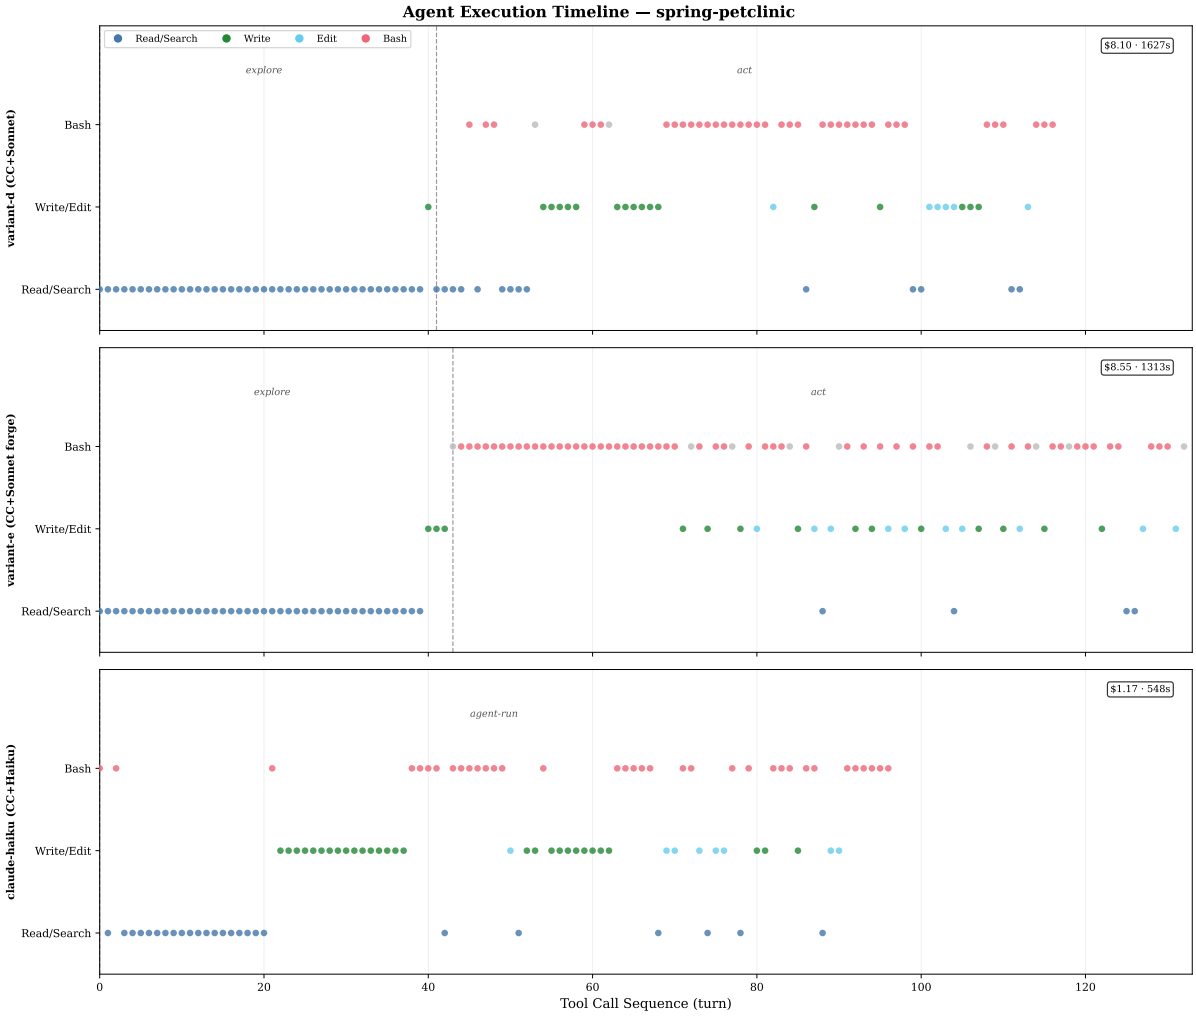

Figure 7: Agent execution timeline — Pet Clinic. Each dot is a tool call. Shared X axis shows relative effort across variants.

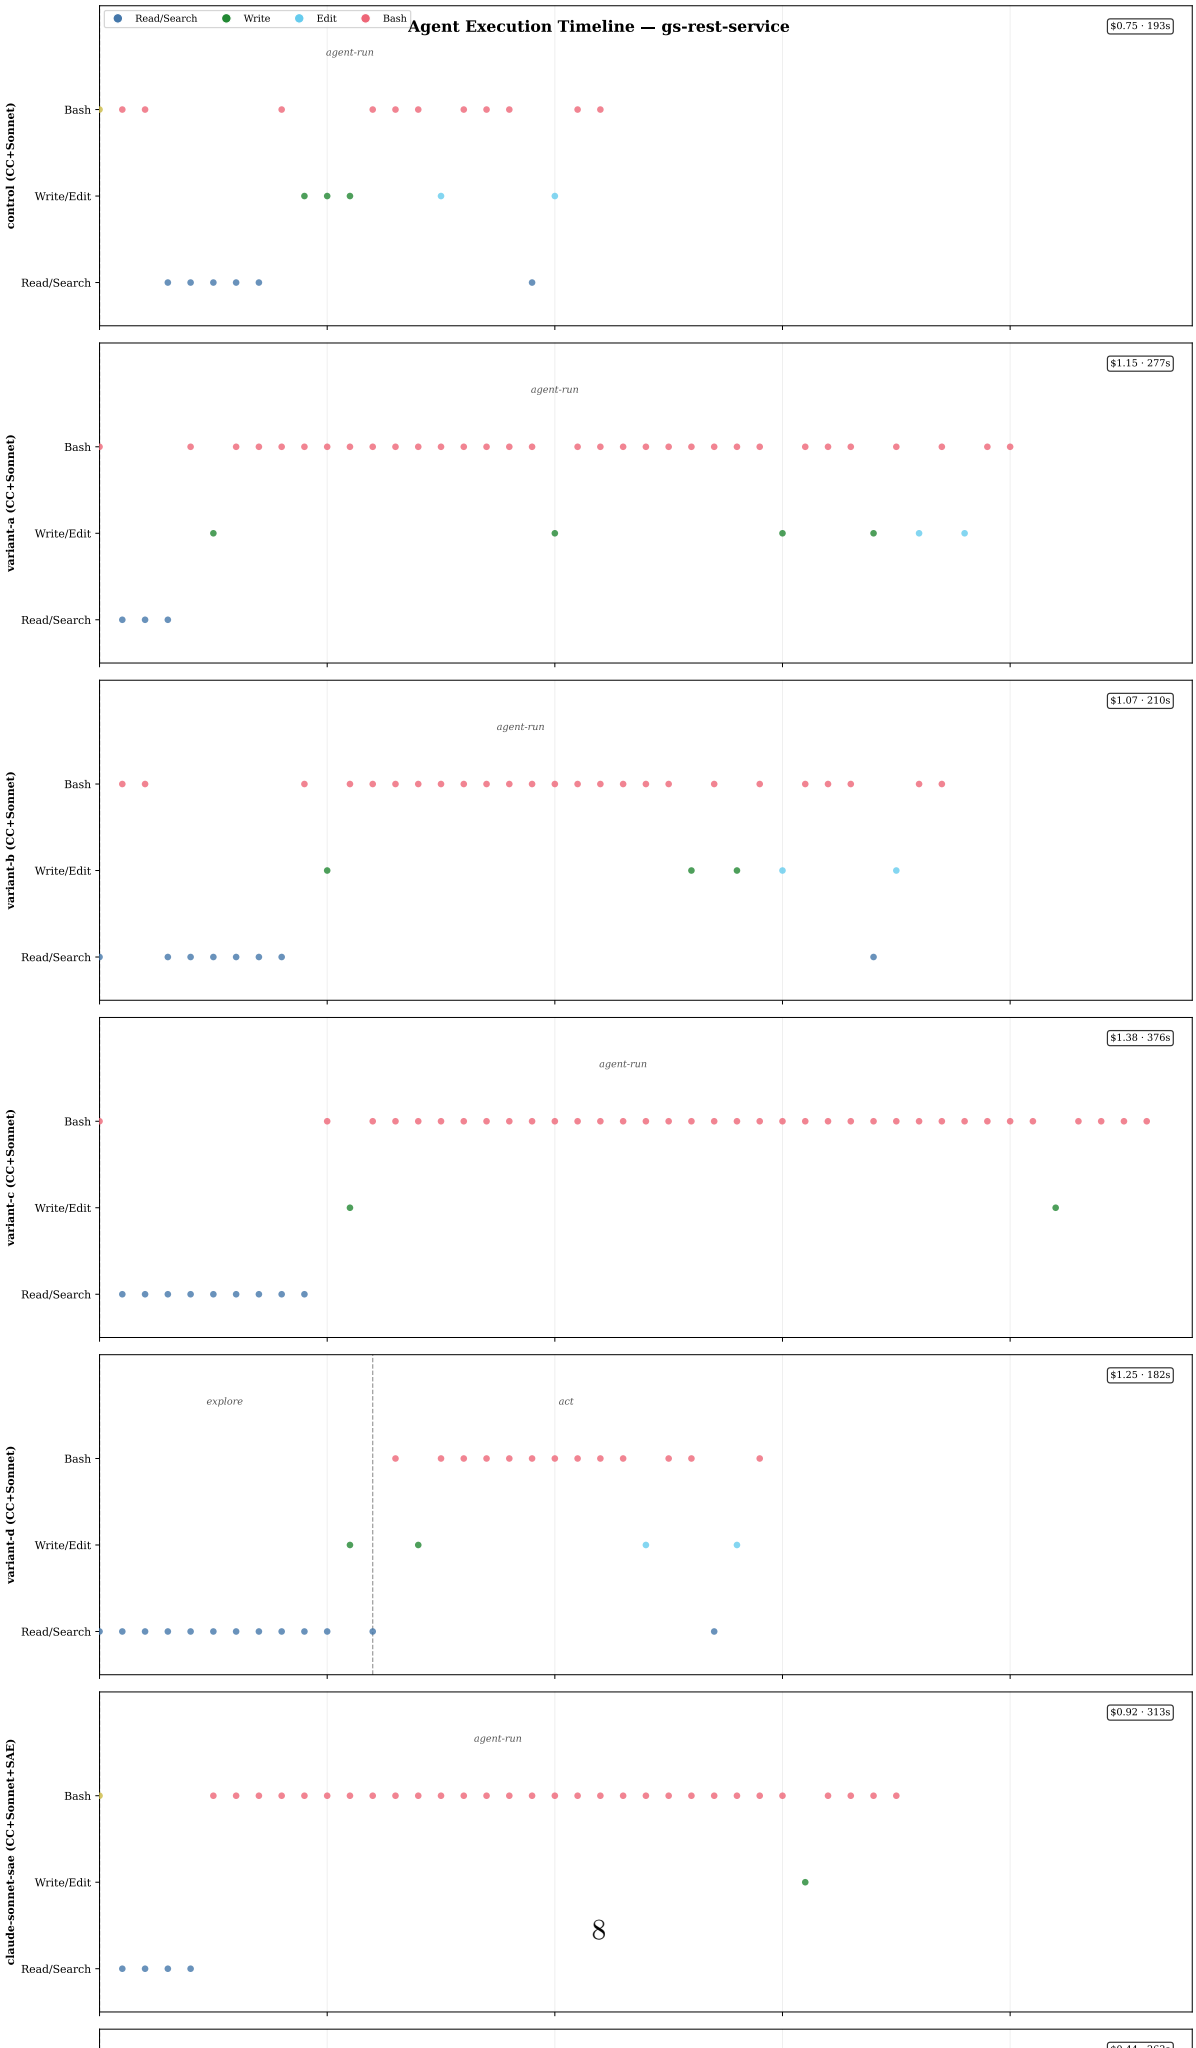
Agent Execution Timeline — gs-rest-service
Read/Search
Write
Edit
Bash
control (CC+Sonnet)
$0.75 · 193s
agent-run
variant-a (CC+Sonnet)
$1.15 · 277s
agent-run
variant-b (CC+Sonnet)
$1.07 · 210s
agent-run
variant-c (CC+Sonnet)
$1.38 · 376s
agent-run
variant-d (CC+Sonnet)
$1.25 · 182s
explore
act
claude-sonnet-sae (CC+Sonnet+SAE)
$0.92 · 313s
agent-run
$0.44 · 263s
Bash
Write/Edit
Read/Search

Table 5: T3 practice adherence per item (higher is better)

Item	CC:Sonnet (naive)	CC:Sonnet (hardened)	CC:Sonnet (+3 KB)	CC:Sonnet (+fu
gs-accessing-data-jpa	0.60	0.70	0.65	0.68
gs-messaging-stomp-websocket	0.50	0.60	0.50	0.50
gs-reactive-rest-service	0.50	0.93	0.70	0.90
gs-rest-service	0.67	0.88	0.85	0.88
gs-securing-web	0.83	0.88	0.78	0.78
spring-petclinic	—	—	—	—

Table 6: Efficiency breakdown per variant (higher is better, except tokens/duration)

Agent:Model	Build Err	Cost Eff	Recovery	Composite	Avg Out Tok	Avg Think Tok	Avg Durat
CC:Sonnet (naive)	0.9	0.82	0.9	0.88	11,086	2,032	221s
CC:Sonnet (hardened)	0.97	0.83	0.97	0.94	9,869	1,558	216s
CC:Sonnet (+3 KB)	0.85	0.8	0.85	0.84	11,894	2,266	237s
CC:Sonnet (+full KB)	0.88	0.79	0.88	0.85	11,203	1,985	334s
CC:Sonnet (two-phase)	0.65	0.55	0.75	0.65	26,436	4,866	508s
CC:Sonnet (forge)	0	0	0.62	0.17	60,528	6,854	1313s
CC:Sonnet (+SAE)	0.88	0.82	0.88	0.86	8,747	1,462	313s
CC:Haiku	0.67	0.9	0.81	0.77	25,472	3,752	294s
CC:Haiku (+SAE)	0	0.9	0.62	0.41	25,680	3,676	332s
Loopy:Sonnet	1	0.46	1	0.86	10,237	0	227s
Loopy:Haiku	1	0.49	1	0.86	26,448	0	330s
Loopy:Qwen3	1	0	1	0.73	24,486	0	1419s

5 Per-Item Breakdown

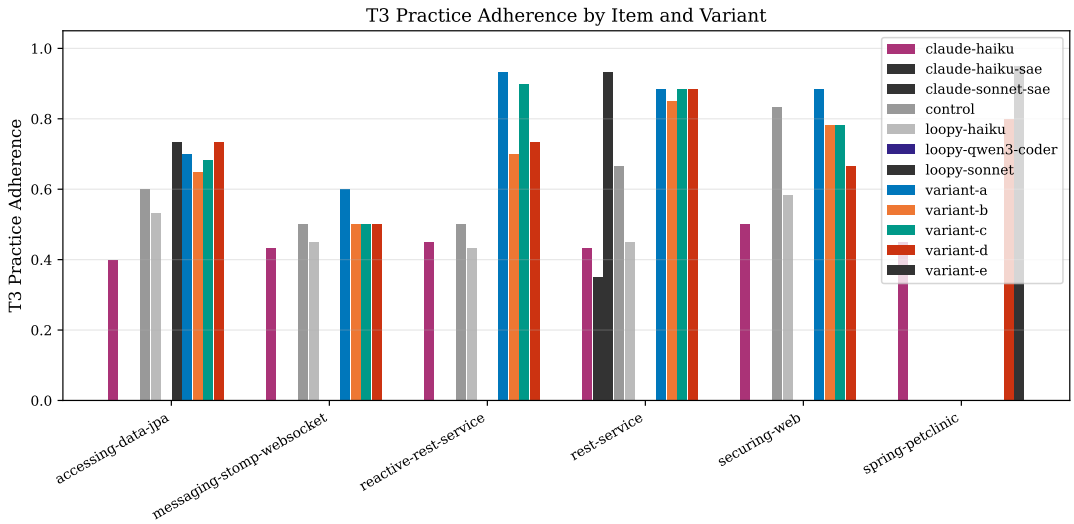

Figure 9: T3 practice adherence by item and variant.

6 Efficiency

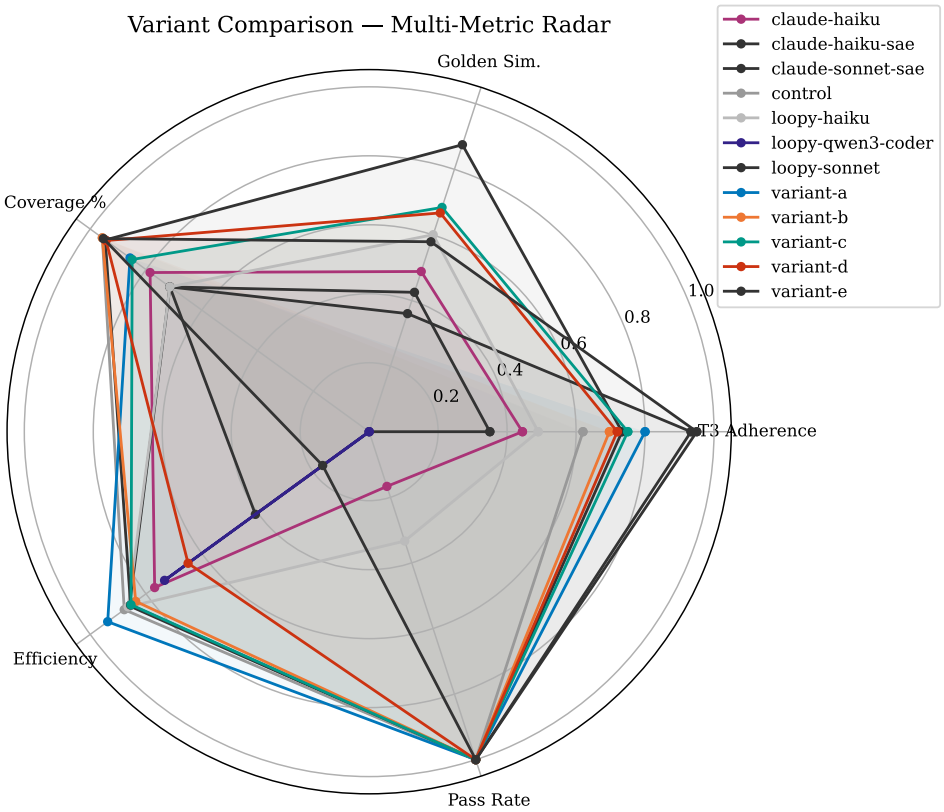

Figure 10: Multi-metric radar chart across all variants.

7 Summary

1. **Prompt engineering is the biggest lever — and the cheapest.** T3: 0.62 → 0.81. +29% quality, same model.

2. **Coverage lies.** 94% coverage, T3=0.45. Without the jury you'd ship bad tests and never know.

3. **There is a model floor.** Below it, no amount of tooling, knowledge, or prompting helps. Haiku, Qwen3 — can't do it.

4. **Model is the quality gate; framework is the cost gate.** Loopy:Sonnet (T3=0.73) ≈ CC:Sonnet (T3=0.70) — same quality, 3× cost. The model determines *whether it works*; the framework determines *what it costs*.

5. **Context compaction is the hidden cost multiplier.** Loopy without compaction: 18M input tokens. With compaction (threshold 0.3): 854K. 21× reduction. Compaction tuning matters.

6. **Knowledge amplifies capability — it cannot create it.** Sonnet + SAE: T3=0.93. Haiku + SAE: T3=0.35. Same pre-computed analysis, opposite results.

7. **Partial knowledge is worse than none.** Giving 3 files triggers more exploration than giving the full KB or nothing at all.

8. **Structured execution shines on hard tasks.** Pet Clinic with plan/act: T3=0.95. Two-phase adds 2× cost on simple targets but pays off on complex ones.